

Vector Differential Calculus in Statistics & Machine Learning

July 30, 2023

1 Motivation-Gaussian regression model

Let $(x_i, y_i), 1 \leq i \leq n$, be a set of measurements on two variables x and y , and consider the problem of fitting a line $y = \beta_0 + \beta_1 x$ to the data. The homoscedastic Gaussian regression model is

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}, \quad \boldsymbol{\varepsilon} \sim N(\mathbf{0}, \sigma^2 \mathbf{I}) \quad (1)$$

where

$$\mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix}, \quad \mathbf{X} = \begin{bmatrix} 1 & x_1 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix} \quad \text{and} \quad \boldsymbol{\beta} = \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix} \quad (2)$$

Assume that σ^2 is known.

1.1 scalar differential calculus approach

The residual sum of squares is

$$\ell(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2 \quad (3)$$

and

$$\frac{\partial \ell(\boldsymbol{\beta})}{\partial \beta_0} = 2 \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i), \quad (4)$$

$$\frac{\partial \ell(\boldsymbol{\beta})}{\partial \beta_1} = 2 \sum_{i=1}^n x_i (y_i - \beta_0 - \beta_1 x_i). \quad (5)$$

Setting this to zero, we obtain

$$\begin{aligned}\hat{\beta}_0 &= \bar{y} - \hat{\beta}_1 \bar{x} \\ \hat{\beta}_1 &= \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) / \sum_{i=1}^n (x_i - \bar{x})^2.\end{aligned}$$

1.2 Vector differential calculus approach

Recall the homoscedastic Gaussian regression model is

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}, \quad \boldsymbol{\varepsilon} \sim N(\mathbf{0}, \sigma^2 \mathbf{I}) \quad (6)$$

where

$$\mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix}, \quad \mathbf{X} = \begin{bmatrix} 1 & x_1 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix} \quad \text{and} \quad \boldsymbol{\beta} = \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix} \quad (7)$$

where we assume that σ^2 is known.

The residual sum of squares is

$$\ell(\boldsymbol{\beta}) = (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^\top (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) \quad (8)$$

Note that (equation 37, The matrix cookbook):

$$\partial(UV) = (\partial U)V + U(\partial V) \quad (9)$$

and

$$\partial U^T = (\partial U)^T \quad (10)$$

Therefore

$$\ell(\boldsymbol{\beta}) = (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^\top (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) \quad (11)$$

$$\implies d\ell(\boldsymbol{\beta}) = \{d(\mathbf{y} - \mathbf{X}\boldsymbol{\beta})\}^\top (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) + (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^\top \{d(\mathbf{y} - \mathbf{X}\boldsymbol{\beta})\} \quad (12)$$

$$= -(\mathbf{X}d\boldsymbol{\beta})^\top (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) - (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^\top \mathbf{X}d\boldsymbol{\beta} \quad (13)$$

$$= (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^\top \mathbf{X}d\boldsymbol{\beta} \quad (14)$$

$$\implies D\ell(\boldsymbol{\beta}) = (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^\top \mathbf{X} = \mathbf{0} \text{ iff } (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^\top \mathbf{X} = \mathbf{0} \quad (15)$$

$$\implies \boldsymbol{\beta} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}. \quad (16)$$

2 Derivatives

2.1 Scalar Case

For the scalar case: given a function $f : \mathbb{R} \rightarrow \mathbb{R}$, then:

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h} \quad (17)$$

$f'(x)$ tells us how much the function f changes as the input x changes by a small amount ε :

$$f(x + \varepsilon) \approx f(x) + \varepsilon f'(x) \quad (18)$$

The **chain rule** tells us how to compute the derivative of the composition of functions. In the scalar case suppose that $f, g : \mathbb{R} \rightarrow \mathbb{R}$ and $y = f(x), z = g(y)$; then we can also write $z = (g \circ f)(x)$. The chain rule tells us that

$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial g(f(x))}{\partial f(x)} \frac{\partial f(x)}{\partial x} \quad (19)$$

Combining these two rules lets us compute the effect of x on z : if x changes by Δx then y will change by $\frac{\partial y}{\partial x} \Delta x$, so we have $\Delta y = \frac{\partial y}{\partial x} \Delta x$. If y changes by Δy then z will change by $\frac{\partial z}{\partial y} \Delta y = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} \Delta x$ which is exactly what the chain rule tells us.

2.2 Gradient

This same intuition carries over into the vector case. Now suppose that $f : \mathbb{R}^N \rightarrow \mathbb{R}$ takes a vector as input and produces a scalar. The derivative of f at the point $x \in \mathbb{R}^N$ is now called the gradient, and it is defined as:

$$\nabla_x f(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{\|h\|} \quad (20)$$

$\nabla_x f(x) \in \mathbb{R}^N$ is a vector, where the i th coordinate of $\frac{\partial y}{\partial x}$ tells us the approximate amount by which y will change if we move x along the i th coordinate axis. We can also view the gradient $\frac{\partial y}{\partial x}$ as a vector of partial derivatives:

$$\frac{\partial y}{\partial x} = \left(\frac{\partial y}{\partial x_1}, \frac{\partial y}{\partial x_2}, \dots, \frac{\partial y}{\partial x_N} \right)$$

where x_i is the i th coordinate of the vector x , which is a scalar, so each partial derivative $\frac{\partial y}{\partial x_i}$ is also a scalar.

Example 1.

$$y(x) = x_1^2 + x_2 + x_2^3 \quad (21)$$

Then

$$\frac{\partial y}{\partial x} = \left(\frac{\partial y}{\partial x_1}, \frac{\partial y}{\partial x_2} \right) = (2x_1, 1 + 3x_2^2) \quad (22)$$

Example 2.

$$y(x) = x_1 + x_2 + x_3^3 \quad (23)$$

Then

$$\frac{\partial y}{\partial x} = \left(\frac{\partial y}{\partial x_1}, \frac{\partial y}{\partial x_2}, \frac{\partial y}{\partial x_3} \right) = (1, 1, 3x_3^2) \quad (24)$$

2.3 Jacobian

Suppose that $f : \mathbb{R}^N \rightarrow \mathbb{R}^M$ takes a vector as input and produces a vector as output. Then the derivative of f at a point x , also called the Jacobian, is the $M \times N$ matrix of partial derivatives. Let $y = f(x)$, then:

$$\frac{\partial y}{\partial x} = \begin{pmatrix} \frac{\partial y_1}{\partial x} \\ \vdots \\ \frac{\partial y_M}{\partial x} \end{pmatrix} = \begin{pmatrix} \frac{\partial y_1}{\partial x_1} & \cdots & \frac{\partial y_1}{\partial x_N} \\ \vdots & \ddots & \vdots \\ \frac{\partial y_M}{\partial x_1} & \cdots & \frac{\partial y_M}{\partial x_N} \end{pmatrix} \quad (25)$$

Example 1.

$$y = (y_1, y_2) = f(x) = (x_1^2 + x_2 + x_2^3, x_1 + x_2 + x_3^3) \quad (26)$$

Then

$$\frac{\partial y_1}{\partial x} = \left(\frac{\partial y_1}{\partial x_1}, \frac{\partial y_1}{\partial x_2}, \frac{\partial y_1}{\partial x_3} \right) = (2x_1, 1 + 3x_2^2, 0) \quad (27)$$

and

$$\frac{\partial y_2}{\partial x} = \left(\frac{\partial y_2}{\partial x_1}, \frac{\partial y_2}{\partial x_2}, \frac{\partial y_2}{\partial x_3} \right) = (1, 1, 3x_3^2) \quad (28)$$

Therefore,

$$\frac{\partial y}{\partial x} = \begin{pmatrix} \frac{\partial y_1}{\partial x} \\ \frac{\partial y_2}{\partial x} \end{pmatrix} = \begin{pmatrix} 2x_1 & 1 + 3x_2^2 & 0 \\ 1 & 1 & 3x_3^2 \end{pmatrix} \quad (29)$$

Multivariate chain rule. The chain rule can be extended to the vector case using Jacobian matrices. Suppose that $f : \mathbb{R}^N \rightarrow \mathbb{R}^M$ and $g : \mathbb{R}^M \rightarrow \mathbb{R}^K$. Let $x \in \mathbb{R}^N$, $y \in \mathbb{R}^M$, and $z \in \mathbb{R}^K$ with $y = f(x)$ and $z = g(y)$, i.e., $z = g(f(x))$. The **multivariate chain rule** says:

$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial g(f(x))}{\partial f(x)} \frac{\partial f(x)}{\partial x} \quad (30)$$

Here, $\frac{\partial z}{\partial y} \in \mathbb{R}^{K \times M}$, $\frac{\partial y}{\partial x} \in \mathbb{R}^{M \times N}$, and $\frac{\partial z}{\partial x} \in \mathbb{R}^{K \times N}$.

Example 3.

$$y(x) = g(f(x)) \quad (31)$$

where $g(z) = z + 1$ and $f(x) = 2x_1 + x_2$. Then by the chain rule

$$\frac{\partial y}{\partial x} = \frac{\partial g(f(x))}{\partial f(x)} \cdot \frac{\partial f(x)}{\partial x} \quad (32)$$

$$= 1 \cdot (2, 1) = (2, 1) \quad (33)$$

2.4 Hessian matrix

Suppose $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is a function taking as input a vector $\mathbf{x} \in \mathbb{R}^n$ and outputting a scalar $f(\mathbf{x}) \in \mathbb{R}$. If all second-order partial derivatives of f exist, then the Hessian matrix is

$$\mathbf{H}_f = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}. \quad (34)$$

That is, the entry of the i th row and the j th column is

$$(\mathbf{H}_f)_{i,j} = \frac{\partial^2 f}{\partial x_i \partial x_j}. \quad (35)$$

2.5 Rules for Differentials

Let \mathbf{u} and \mathbf{v} be vector functions and \mathbf{U} and \mathbf{V} be matrix functions. A will denote a constant matrix and s a scalar function.

2.5.1 Rules for Scalar Functions

$$\begin{aligned} du^\alpha &= \alpha u^{\alpha-1} du, \\ d \log u &= u^{-1} du, \\ de^u &= e^u du. \end{aligned}$$

2.5.2 Rules Involving Linear Functions

$$\begin{aligned} d(\mathbf{AU}) &= \mathbf{A}d\mathbf{U}, \\ d(\mathbf{U} + \mathbf{V}) &= d\mathbf{U} + d\mathbf{V}, \\ d \operatorname{diag}(\mathbf{u}) &= \operatorname{diag}(d\mathbf{u}), \\ d\mathbf{U}^\top &= (d\mathbf{U})^\top, \\ d \operatorname{vec} \mathbf{U} &= \operatorname{vec}(d\mathbf{U}), \\ d(\operatorname{tr} \mathbf{U}) &= \operatorname{tr}(d\mathbf{U}), \\ d(E\mathbf{U}) &= E(d\mathbf{U}). \end{aligned}$$

2.5.3 Rules for Determinant and Matrix Inverse

$$\begin{aligned} d|\mathbf{U}| &= |\mathbf{U}| \operatorname{tr}(\mathbf{U}^{-1} d\mathbf{U}), \\ d\mathbf{U}^{-1} &= -\mathbf{U}^{-1}(d\mathbf{U})\mathbf{U}^{-1}. \end{aligned}$$

2.5.4 Rules Involving Quadratic Forms

$$\begin{aligned} d\mathbf{u}^\top \mathbf{A} \mathbf{u} &= \mathbf{u}^\top (\mathbf{A} + \mathbf{A}^\top) d\mathbf{u}, \\ d\mathbf{u}^\top \mathbf{A} \mathbf{u} &= 2\mathbf{u}^\top \mathbf{A} d\mathbf{u}, \quad \mathbf{A} \text{ symmetric.} \end{aligned}$$

3 Maximum Likelihood Estimate for parameters of a multivariate Gaussian Distribution

Given a set of i.i.d. data $X = \{x_1, \dots, x_N\}$ drawn from $\mathcal{N}(x; \mu, \Sigma)$, we want to estimate (μ, Σ) by MLE. The log-likelihood function is

$$\ln p(X | \mu, \Sigma) = -\frac{N}{2} \ln |\Sigma| - \frac{1}{2} \sum_{n=1}^N (x_n - \mu)^\top \Sigma^{-1} (x_n - \mu) + \text{const} \quad (36)$$

Taking its derivative w.r.t. μ and setting it to zero we have

$$\hat{\mu} = \frac{1}{N} \sum_{n=1}^N x_n \quad (37)$$

Rewrite the log-likelihood using $\text{Trace}(\text{constant}) = \text{constant}$, and $\text{Trace}(AB) = \text{Trace}(BA)$,

$$\ln p(X | \mu, \Sigma) = -\frac{N}{2} \ln |\Sigma| - \frac{1}{2} \sum_{n=1}^N (x_n - \mu)^\top \Sigma^{-1} (x_n - \mu) + \text{const} \quad (38)$$

$$\propto -\frac{N}{2} \ln |\Sigma| - \frac{1}{2} \sum_{n=1}^N \text{Trace} \left(\Sigma^{-1} (x_n - \mu) (x_n - \mu)^\top \right) \quad (39)$$

$$= -\frac{N}{2} \ln |\Sigma| - \frac{1}{2} \text{Trace} \left(\Sigma^{-1} \sum_{n=1}^N \left[(x_n - \mu) (x_n - \mu)^\top \right] \right) \quad (40)$$

Taking the derivative w.r.t. Σ^{-1} , and using 1) $\frac{\partial}{\partial A} \log |A| = A^{-T}$; 2) $\frac{\partial}{\partial A} \text{Tr}[AB] = \frac{\partial}{\partial A} \text{Tr}[BA] = B^T$, we obtain

$$\hat{\Sigma} = \frac{1}{N} \sum_{n=1}^N (x_n - \hat{\mu}) (x_n - \hat{\mu})^\top. \quad (41)$$

4 Generalized Linear Models

Let \mathbf{y} be a vector of responses and \mathbf{X} be a corresponding design matrix. The one-parameter exponential family model, with canonical link, is characterized by the joint density

$$f(\mathbf{y}; \beta) = \exp \left\{ \mathbf{y}^\top (\mathbf{X}\beta) - \mathbf{1}^\top b(\mathbf{X}\beta) + \mathbf{1}^\top c(\mathbf{y}) \right\} \quad (42)$$

where β is the vector of coefficients. For example, $b(x) = \log(1 + e^x)$ corresponds to binary regression with a logit link function.

The log-likelihood of β is

$$\ell(\beta) = \mathbf{y}^\top \mathbf{X}\beta - \mathbf{1}^\top b(\mathbf{X}\beta) + \mathbf{1}^\top c(\mathbf{y}) \quad (43)$$

$$\implies d\ell(\beta) = \mathbf{y}^\top \mathbf{X}d\beta - \mathbf{1}^\top db(\mathbf{X}\beta) \quad (44)$$

$$= \mathbf{y}^\top \mathbf{X}d\beta - \mathbf{1}^\top \text{diag}\{b'(\mathbf{X}\beta)\} d(\mathbf{X}\beta) \quad (45)$$

$$= \mathbf{y}^\top \mathbf{X}d\beta - b'(\mathbf{X}\beta)^\top \mathbf{X}d\beta \quad (46)$$

$$= \{\mathbf{y} - b'(\mathbf{X}\beta)\}^\top \mathbf{X}d\beta. \quad (47)$$

Hence,

$$D\ell(\beta) = \{\mathbf{y} - b'(\mathbf{X}\beta)\}^\top \mathbf{X}$$

Also,

$$d^2\ell(\beta) = d\{\mathbf{y} - b'(\mathbf{X}\beta)\}^\top \mathbf{X}d\beta \quad (48)$$

$$= -\{\text{diag}\{b''(\mathbf{X}\beta)\} \mathbf{X}d\beta\}^\top \mathbf{X}d\beta \quad (49)$$

$$= (d\beta)^\top \mathbf{X}^\top [-\text{diag}\{b''(\mathbf{X}\beta)\}] \mathbf{X}(d\beta) \quad (50)$$

which leads to

$$H\ell(\beta) = -\mathbf{X}^\top \text{diag}\{b''(\mathbf{X}\beta)\} \mathbf{X} \quad (51)$$

5 Extra examples and reading

- [Vector Differential Calculus in Statistics](#)
- [Matrix handbook for multivariate statistics and Machine Learning](#)
- [The Matrix Cookbook](#)